

Why not just buy a CRM?

A CRM is a **system of record** — a filing cabinet someone has to fill in by hand. What we built is a **system of intelligence**: it finds the leads, verifies the contacts, writes the record, and turns every meeting into knowledge that compounds. The real choice isn't “CRM vs. this.” It's **a database someone must feed** vs. **a loop that feeds itself**.

One closed loop, across two purpose-built systems

The RocketReach lead engine and the Faragopedia knowledge graph aren't two tools bolted together — they form a cycle. Research flows into outreach; outreach flows into institutional memory; and that memory calibrates the next round of research. A CRM sits at one point on this loop (storage) and leaves every other step to a human.

LEAD ENGINE (ROCKETREACH · SHEETS)

KNOWLEDGE GRAPH (FARAGOPEDIA)

01 · AUTO

Prospect

AI drafts fresh target companies every Monday, tuned to real client fit.



02 · AUTO

Enrich

RocketReach verifies email, title & location from the LinkedIn profile.



03 · HUMAN

Outreach

Pre-filled rows land in the sheet. Siobahn already works in. She just picks who to pursue.

04 · TRIGGER

Meeting booked

Marking a win auto-syncs the record into the knowledge base.



05 · AUTO

Knowledge graph

A linked wiki page is written & cross-referenced — people, companies, context.

← COMPOUNDS BACK

The step a CRM can't do: accumulated intelligence about who we know, why they mattered, and what worked feeds straight back into next week's targeting. The system gets sharper every cycle — on its own.

Three things a subscription CRM structurally will not do

01

It finds who to approach — not just their email

Modern CRMs can enrich a contact from a name and company — but only as a paid add-on, and only once you've decided who that contact is. Ours does the deciding: it proposes fit-calibrated targets each week, verifies them, and writes them in.

Per contact: ~30 min of hunting → 3-5 sec

02

Knowledge compounds instead of sitting in fields

A CRM stores flat records. Faragopedia builds a self-maintaining knowledge graph — every meeting enriches an interlinked web of people, companies and context, with the AI keeping links and summaries current.

Output = institutional memory, not a contact list

03

It lives where the work already happens

No new tool to adopt, no migration, no seats to train. Prospects surface in the Google Sheet Siobahn already uses; intelligence surfaces in a wiki anyone on the team can read.

Adoption cost: zero new software

Same job, two very different machines

	OFF-THE-SHELF CRM	FARAGO BD SYSTEM
FINDING LEADS	You research and type them in	AI proposes calibrated targets weekly
CONTACT DATA	Add a paid enrichment add-on, then map fields	RocketReach-verified, written in automatically
KEEPING IT CURRENT	Manual data entry & hygiene	AI maintains pages, links & contradictions
KNOWLEDGE MODEL	Flat records in fixed fields	Interlinked knowledge graph that compounds
WHERE YOU WORK	A separate app the team must adopt	The spreadsheet & wiki already in use
OWNERSHIP	Rented; intelligence leaves when you cancel	Portable markdown files we own outright

The bottleneck was never the software. It was the hours.

A CRM — even with a paid enrichment add-on — still runs on human time: someone decides who to chase, hunts down the email and title, keys it in, and later digs back through records to recover the context. Automating that research is where the hours come back.

RESEARCH & ENRICHMENT · PER CONTACT

BY HAND CRM + manual research		~20-30 min
AUTOMATED Farago lead engine		3-5 sec

10-30x
MORE THROUGHPUT

Freed from the research grind, one person works through **10-30x the prospects** — a week of hunting collapses into an afternoon. And it **scales**: a single run handles 5 contacts or 500 at essentially the same time cost, so more volume no longer means another hire. Retrieval is the same story — context that takes minutes to dig out of a CRM's notes is **one search** here.

The honest accounting

What the off-the-shelf route actually costs

- + CRM **seats** — **\$12–\$175+** per user / month, billed annually
- + A **data-enrichment** add-on to auto-fill contacts — **~\$49–\$99/mo**, or **~\$0.20+** per contact beyond plan credits
- + **Integration middleware** (Zapier / Make) to wire it together
- + **Onboarding & setup** — **\$1,500–\$3,500** one-time, plus certified consultants at **\$100–\$300/hr**
- + ... and you'd **still do the prospecting by hand**

Typical 2026 market pricing. The sticker price is only the first line — seats, enrichment, middleware and setup stack into the real total, and none of it decides who to approach for you.

Where a CRM genuinely wins

- ✓ Mature **deal-pipeline** mechanics & forecasting
- ✓ Native **email & calendar** logging
- ✓ **Multi-user** permissions and large-team workflows
- ✓ Out-of-the-box **dashboards**, mobile apps, vendor support

We deliberately don't compete on sales-funnel mechanics for a big team. This is built for a focused BD workflow — and if we ever need pipeline mechanics, **this loop can feed a CRM rather than be replaced by one.**

— BOTTOM LINE

A CRM would give us a **differently-organised place to store contacts we'd still have to find ourselves.** This gives us the contacts, the context, and a knowledge base that compounds — inside the tools the team already uses. It isn't a cheaper CRM. It's a different, and for our purpose better, machine.